

Aditya Dasika

LLM Inference & Serving Platform · LLM Evaluation, Monitoring & Guardrails · Multi-Tenant Agent Orchestration

akhil.dasika47@gmail.com | [+91-96184-33030](tel:+91-96184-33030) | [Portfolio](#) | [GitHub](#) | [LinkedIn](#)

EXPERIENCE

AI/LLM Engineer II

Phenom People

April 2025 – Present

Hyderabad, Telangana

- Ran the self-hosted LLM inference service as a deployable platform — Helm-packaged releases, custom chat templates per model, reasoning-model flags, and **OpenTelemetry** instrumentation feeding a metrics dashboard so per-model latency and throughput were visible in production.
- Built an **LLM evaluation framework** (pass@k, LLM-as-judge rubric) gating every release across all 12 nodes, instrumented with **LiteLLM** + **Langfuse** for per-node cost, latency and full-trace auditability — continuously monitoring agent behaviour and curbing hallucinations.
- Fine-tuned and served Qwen models via **QLoRA** (Unsloth, NF4) on **vLLM** / AWS g5.xlarge (A10G), sustaining 10.98 RPS at 98.6% success under load at roughly half the inference cost of hosted GPT-4.1.
- Hardened the fine-tuned Qwen deployment for production serving — a custom Jinja2 chat template, required tool-choice and temperature pinned at 0.1 for deterministic tool calls — then held it to a golden baseline of 1.0s median and **1.4s** P95 latency with zero failures across **5,510** requests, taking the service's test coverage from **7% to 78%**.
- Architected **chatbotCXAgent**, a production multi-agent system on **LangGraph** (12 orchestrated nodes) spanning intent recognition, planning, dialogue-state management and tool routing — with **Graph RAG** retrieval over enterprise knowledge and document repositories via FAISS vector search and a Neo4j graph for multi-hop, policy-aware grounding.
- Engineered tool-calling agents that execute governed actions — database queries and structured extraction from unstructured documents — under **XGrammar constrained decoding** and schema-validated JSON guardrails, raising tool-routing accuracy to **91.9%** and eliminating malformed tool calls in live traffic.
- Built the fraud-risk configuration service greenfield in **FastAPI** on **MongoDB** — signal registry, feature groups, audience and scope definitions and tenant-scoped overrides — exposed through an internal resolve endpoint that **Temporal** workflows call at workflow start to decide which fraud signals apply to a candidate and job.
- Building the detection and decisioning layer of an enterprise **hiring-integrity / fraud-risk platform** — an **LLM orchestration engine** that investigates flagged candidate activity, correlates evidence across signals and produces reviewer-facing assessments with supporting rationale — backed by a **policy-driven rules engine**, per-tenant risk configuration and audited review workflows across the signal lifecycle.
- Owned tenant provisioning and delivery for the platform — provision and unprovision APIs with region-scoped tenant listing, containerised deployment via **Docker** and **Helm** across staging environments, and production support debugging by querying service logs in **ClickHouse** to resolve customer tickets.
- Built conversation analytics APIs covering both live-agent and WhatsApp channels, raised the service's automated testing coverage from **63% to over 90%** while clearing static-analysis code smells, and migrated the service to **FastAPI**.

Full Stack Engineer

Blueleaves Farms

May 2022 – February 2025

Hyderabad, Telangana

- Owned end-to-end delivery of production systems (POS, CRM, inventory, procurement, e-commerce) on Node.js, Firebase Cloud Functions and REST APIs — with role-based access control and CI/CD — integrating multiple internal systems and cutting manual operational errors by roughly 90%.

TECHNICAL SKILLS

AI Platform & Serving: vLLM, Helm, Docker, Production MLOps, Model Routing, LLM Fine-Tuning (QLoRA, SFT), Unsloth, HuggingFace PEFT, Ollama, AWS (EC2/S3/SageMaker)

Evaluation, Observability & Governance: LLM Evaluation (LLM-as-judge, pass@k), OpenTelemetry, Langfuse, LiteLLM, Model Monitoring, Constrained Decoding (XGrammar), Schema Validation, Hallucination Mitigation, Role-Based Access Control

GenAI & Agentic: Agentic / Multi-Agent Systems, LangGraph, LangChain, Tool/Function Calling, Prompt Engineering, Context-Window Management, RAG & Graph RAG, Embeddings / Vector Search (FAISS)

Backend, APIs & Delivery: Python, TypeScript, JavaScript, SQL, FastAPI, Node.js, REST APIs, Microservices, Temporal, Kafka, GitHub Actions, CI/CD, Linux

Databases: PostgreSQL, MongoDB, Neo4j, ClickHouse, Redis, Firebase

EDUCATION

Andhra Loyola Institute of Engineering and Technology

Bachelor of Technology in Mechanical Engineering

Vijayawada, Andhra Pradesh

August 2018 – May 2022