

Aditya Dasika

AI Platform Engineer · Agent Orchestration on LangGraph · OpenTelemetry Observability & LLM Evaluation · Docker/Helm Deployment

akhil.dasika47@gmail.com | [+91-96184-33030](tel:+91-96184-33030) | [Portfolio](#) | [GitHub](#) | [LinkedIn](#)

EXPERIENCE

AI/LLM Engineer II

April 2025 – Present

Phenom People

Hyderabad, Telangana

- Ran the self-hosted LLM inference service as a deployable platform — Helm-packaged releases, custom chat templates per model and reasoning-model flags — instrumented with **OpenTelemetry** feeding a metrics dashboard, so per-model latency and throughput were visible in production.
- Owned tenant provisioning and delivery for the platform — provision and unprovision APIs with region-scoped tenant listing and environment configuration, containerised deployments via **Docker** and **Helm** across staging environments — and ran production support, troubleshooting by querying service logs in **ClickHouse** to resolve customer tickets.
- Built an **LLM evaluation framework** (pass@k, LLM-as-judge rubric) gating every release across all 12 nodes, instrumented with LiteLLM and Langfuse for per-node cost, latency and full-trace auditability — continuously monitoring agent behaviour and curbing hallucinations.
- Architected **chatbotCXAgent**, a production multi-agent system on **LangGraph** whose agent orchestration spans 12 nodes — intent recognition, planning, dialogue-state management and tool routing — with **Graph RAG** retrieval over enterprise knowledge and document repositories via FAISS vector search and a Neo4j graph for multi-hop, policy-aware grounding.
- Engineered tool-calling agents that execute governed actions — database queries and structured extraction from unstructured documents — under **XGrammar constrained decoding** and schema-validated JSON guardrails, raising tool-routing accuracy to **91.9%** and eliminating malformed tool calls in live traffic.
- Fine-tuned and served Qwen models via **QLoRA** on **vLLM** / AWS g5.xlarge (A10G) at roughly half the inference cost of hosted GPT-4.1, sustaining 10.98 RPS at 98.6% success under load, then hardened the deployment for production — temperature pinned at 0.1 for deterministic tool calls, load testing holding a golden baseline of 1.0s median and **1.4s** P95 latency with zero failures across **5,510** requests, and test coverage from **7% to 78%**.
- Built the fraud-risk configuration service greenfield in **FastAPI** on **MongoDB** — signal registry, feature groups, audience and scope definitions, and tenant-scoped overrides — exposed through an internal resolve endpoint that **Temporal** workflows call at workflow start to decide which fraud signals apply to a candidate and job.

Full Stack Engineer

May 2022 – February 2025

Blueleaves Farms

Hyderabad, Telangana

- Owned end-to-end delivery of production systems (POS, CRM, inventory, procurement, e-commerce) on Node.js, Firebase Cloud Functions and REST APIs — with role-based access control and CI/CD pipelines — integrating multiple internal systems and cutting manual operational errors by roughly 90%.
- **Sole engineer** on the whole stack — built the CRM and the other internal tools as **React** and **Vite** web applications, owning the interface as well as the services behind it.

PROJECTS

IndexNotes — AI Knowledge Graph for Scientific Papers | *Neo4j, Ollama, PyMuPDF, Python, Docker*

- Built an automated Graph RAG pipeline ingesting research papers (PDF to local-LLM extraction to structured JSON to Neo4j), creating 40+ concepts and 17+ typed relationships, with quantized local models via Ollama for entity and relationship extraction.
- Designed a Neo4j schema enabling multi-hop graph traversal and cross-document retrieval across scientific domains.

TECHNICAL SKILLS

Agentic Frameworks & GenAI: Agentic / Multi-Agent Systems, LangGraph, LangChain, Tool/Function Calling, Prompt Engineering, Model Routing, Context-Window Management, RAG & Graph RAG, LLM Fine-Tuning (QLoRA, SFT)
Observability, Evaluation & Guardrails: OpenTelemetry, Langfuse, LiteLLM, Model Monitoring, LLM Evaluation (LLM-as-judge, pass@k), Constrained Decoding (XGrammar), Schema Validation, Hallucination Mitigation, Role-Based Access Control

Languages: Python, JavaScript, TypeScript, SQL, Bash

Platform, Deployment & Infra: FastAPI, Docker, Helm, Temporal, Kafka, REST APIs, Microservices, AWS (EC2/S3/SageMaker), GitHub Actions, CI/CD, Linux, Node.js

ML & Serving: vLLM, PyTorch, Transformers (BERT/ModernBERT), HuggingFace PEFT, Unsloth, Embeddings / Vector Search (FAISS), Neo4j, Ollama, Production MLOps

Databases: PostgreSQL, MongoDB, ClickHouse, Redis, Neo4j, Firebase

EDUCATION

Andhra Loyola Institute of Engineering and Technology

Bachelor of Technology in Mechanical Engineering

Vijayawada, Andhra Pradesh

August 2018 – May 2022