

Aditya Dasika

akhil.dasika47@gmail.com | +91-96184-33030 | LinkedIn

August 24, 2026

Dear Anandha,

Thank you for reaching out about the Gen AI Developer role in Bengaluru. I build agentic LLM systems end to end, from orchestration down to eval, and shipped one to production this year: chatbotCXAgent, a 12-node LangGraph platform now handling live enterprise workflows at Phenom People.

TCS's push into agentic AI, from MasterCraft's GenAI-driven legacy modernization (50 million lines of COBOL migrated to Java for a financial services client) to the orchestration layer in WisdomNext, is the kind of problem I want to work on: agents doing real, measured work inside large systems. Here's what I'd bring to the Gen AI Developer scope:

- **Agentic orchestration:** Designed and shipped a 12-node LangGraph platform (intent recognition, dialogue state, tool-routing, Graph RAG) to production users
- **Tool-call reliability:** XGrammar constrained decoding, schema-validated outputs, 91.9% tool-routing accuracy and zero malformed calls in production
- **RAG and retrieval:** FAISS-based Graph RAG pipelines, including an independent project turning scientific PDFs into a queryable Neo4j knowledge graph
- **Fine-tuning, serving, and eval:** QLoRA on vLLM at 10.98 RPS / 98.6% success, roughly halving inference cost, gated by a pass@k and LLM-judge framework

I gate my own releases behind eval frameworks rather than shipping on faith, and I own systems past the demo stage, including the reliability and cost work a production agent needs once real traffic hits it. That is the standard I would bring to a GenAI Developer role at TCS.

I look forward to the interview on the 29th.

Kind regards,

Aditya Dasika